

Introduction to MicroBlaze

LECTURE 5

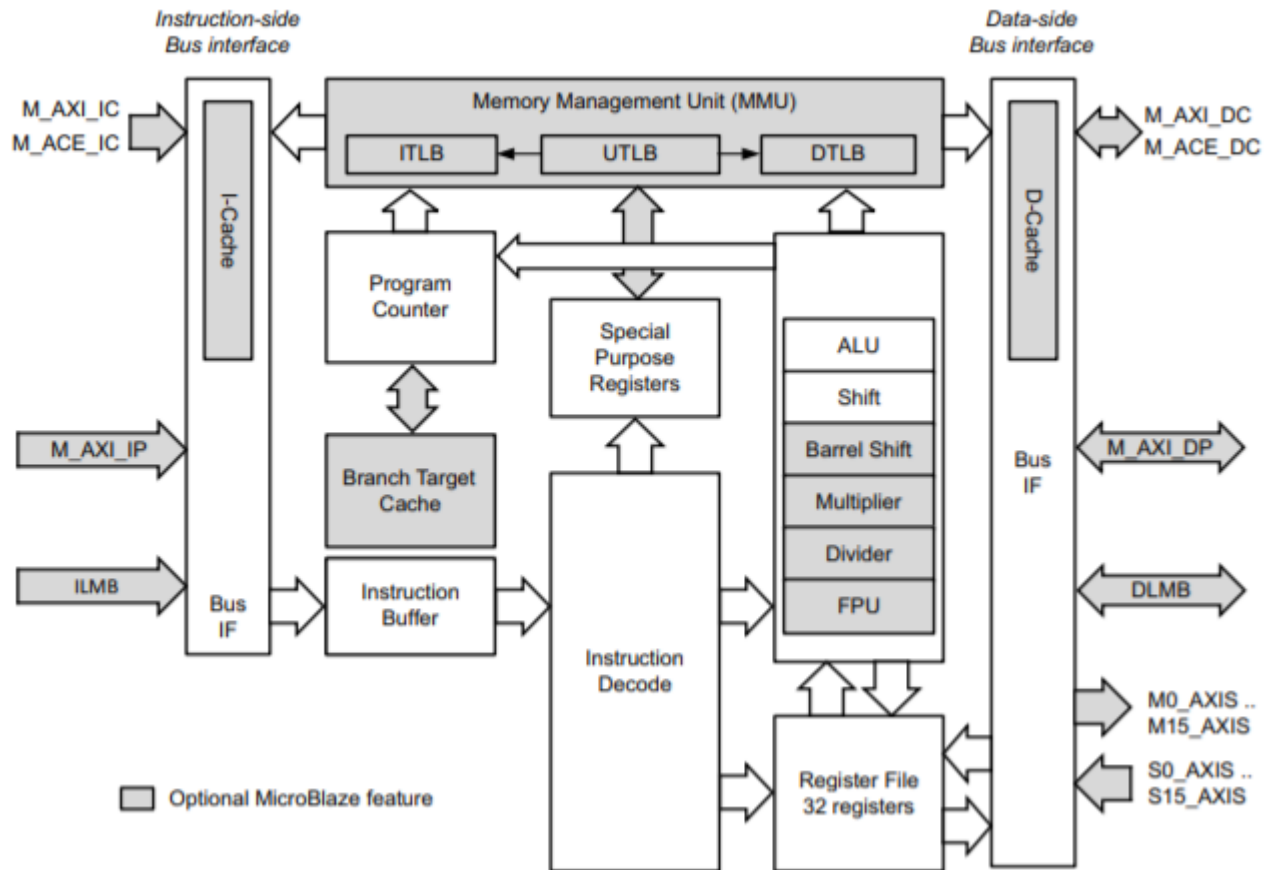
PEDRO SANTOS (SLIDES BY IOULIIA SKLIAROVA)

MicroBlaze Processor

- Soft processor
 - ~1 900-7 000 logic cells (~1 200-4 500 LUTs) - estimates
 - 63 400 LUTs available in Artix-7 XC7A100T
- RISC architecture
 - utilizes a small, highly-optimized set of instructions, rather than a more specialized set of instructions often found in other types of architectures
- 32/64-bit architecture
 - Thirty-two 32-bit or 64-bit general purpose registers
- In production since 2002
- Supported in
 - 7-series/UltraScale/UltraScale+ devices
- Three preset configurations:
 - a simple microcontroller running bare-metal applications (~200 MHz in Artix-7);
 - a real-time processor running FreeRTOS (~170 MHz);
 - an application processor with a memory management unit running Linux (~140 MHz)

MicroBlaze Overview

MicroBlaze Processor Reference Guide - UG984



MicroBlaze Features

- The fixed feature set of the processor includes:
 - Thirty-two 32-bit or 64-bit general purpose registers
 - 32-bit instruction word with three operands and two addressing modes
 - Default 32-bit address bus, extensible to 64 bits
 - Single issue pipeline
- Configurable features
 - Processor pipeline depth
 - Floating-point unit (FPU)
 - Hardware divider
 - Optimizable for Area, Performance or Frequency
 - 64-bit mode
 - ...

MicroBlaze Data Types

- The MicroBlaze processor uses Big-Endian or Little-Endian (default) format to represent data, depending on the selected endianness.
- The hardware supported data types for 32-bit MicroBlaze are word, half word, and byte. With 64-bit MicroBlaze the data types long and double are also available in hardware.

Big-Endian Byte Address	n	n+1	n+2	n+3
Big-Endian Byte Significance	MSByte			LSByte
Big-Endian Byte Order	n	n+1	n+2	n+3
Big-Endian Byte-Reversed Order	n+3	n+2	n+1	n
Little-Endian Byte Address	n+3	n+2	n+1	n
Little-Endian Byte Significance	MSByte			LSByte
Little-Endian Byte Order	n+3	n+2	n+1	n
Little-Endian Byte-Reversed Order	n	n+1	n+2	n+3
Bit Label	0			31
Bit Significance	MSBit			LSBit

MicroBlaze Instruction Summary

- All MicroBlaze instructions are 32 bits and are defined as either Type A or Type B.
- Type A instructions have up to two source register operands and one destination register operand.
- Type B instructions have one source register and a 16-bit immediate operand.
Type B instructions have a single destination register operand.

Table 2-7: MicroBlaze Instruction Set Summary

Type A	0-5	6-10	11-15	16-20	21-31	Semantics
Type B	0-5	6-10	11-15	16-31		
ADD Rd,Ra,Rb	000000	Rd	Ra	Rb	00L00000000	$Rd := Rb + Ra$
RSUB Rd,Ra,Rb	000001	Rd	Ra	Rb	00L00000000	$Rd := Rb + \overline{Ra} + 1$
ADDC Rd,Ra,Rb	000010	Rd	Ra	Rb	00L00000000	$Rd := Rb + Ra + C$
RSUBC Rd,Ra,Rb	000011	Rd	Ra	Rb	00L00000000	$Rd := Rb + \overline{Ra} + C$
ADDK Rd,Ra,Rb	000100	Rd	Ra	Rb	00L00000000	$Rd := Rb + Ra$
RSUBK Rd,Ra,Rb	000101	Rd	Ra	Rb	00L00000000	$Rd := Rb + \overline{Ra} + 1$
CMP Rd,Ra,Rb	000101	Rd	Ra	Rb	00L00000001	$Rd := Rb + \overline{Ra} + 1$ $Rd[0] := 0$ if $(Rb \geq Ra)$ else $Rd[0] := 1$
CMPU Rd,Ra,Rb	000101	Rd	Ra	Rb	00L00000011	$Rd := Rb + \overline{Ra} + 1$ (unsigned) $Rd[0] := 0$ if $(Rb \geq Ra, \text{ unsigned})$ else $Rd[0] := 1$

MicroBlaze Pipeline

MicroBlaze Pipeline

- MicroBlaze instruction execution is pipelined.
- For most instructions, each stage takes one clock cycle to complete.
- Consequently, the number of clock cycles necessary for a specific instruction to complete is equal to the number of pipeline stages, and one instruction is completed on every cycle in the absence of data, control or structural hazards.
 - A data hazard occurs when the result of an instruction is needed by a subsequent instruction. This can result in stalling the pipeline, unless the result can be forwarded to the subsequent instruction. The MicroBlaze GNU Compiler attempts to avoid data hazards by reordering instructions during optimization.
 - A control hazard occurs when a branch is taken, and the next instruction is not immediately available. This results in stalling the pipeline. MicroBlaze provides delay slot branches and the optional branch target cache to reduce the number of stall cycles.
 - A structural hazard occurs for a few instructions that require multiple clock cycles in the execute stage or a later stage to complete. This is achieved by stalling the pipeline.

Three Stage Pipeline

- With the MicroBlaze is optimized for **area**, the pipeline is divided into three stages to minimize hardware cost: Fetch, Decode, and Execute.
- The three stage pipeline does not have any data hazards. Pipeline stalls are caused by control hazards, structural hazards due to multi-cycle instructions, memory accesses using slower memory, instruction fetch from slower memory, or stream accesses.

	cycle1	cycle2	cycle3	cycle4	cycle5	cycle6	cycle7
instruction 1	Fetch	Decode	Execute				
instruction 2		Fetch	Decode	Execute	Execute	Execute	
instruction 3			Fetch	Decode	Stall	Stall	Execute

Five Stage Pipeline

- With the MicroBlaze is optimized for **performance**, the pipeline is divided into five stages to maximize performance: Fetch (IF), Decode (OF), Execute (EX), Access Memory (MEM), and Writeback (WB).
- Pipeline stalls are caused by data hazards, control hazards, structural hazards due to multicycle instructions, memory accesses using slower memory, instruction fetch from slower memory, or stream accesses.

	cycle1	cycle2	cycle3	cycle4	cycle5	cycle6	cycle7	cycle8	cycle9
instruction 1	IF	OF	EX	MEM	WB				
instruction 2		IF	OF	EX	MEM	MEM	MEM	WB	
instruction 3			IF	OF	EX	Stall	Stall	MEM	WB

Eight Stage Pipeline

- With the MicroBlaze is optimized for **frequency**, the pipeline is divided into eight stages to maximize possible frequency: Fetch (IF), Decode (OF), Execute (EX), Access Memory 0 (M0), Access Memory 1 (M1), Access Memory 2 (M2), Access Memory 3 (M3) and Writeback (WB).
- Pipeline stalls are caused by data hazards, control hazards, structural hazards, memory accesses using slower memory, instruction fetch from slower memory, or stream accesses.

	cycle1	cycle2	cycle3	cycle4	cycle5	cycle6	cycle7	cycle8	cycle9	cycle10	cycle11
instruction 1	IF	OF	EX	M0	M1	M2	M3	WB			
instruction 2		IF	OF	EX	M0	M0	M1	M2	M3	WB	
instruction 3			IF	OF	EX	Stall	M0	M1	M2	M3	WB

Memory Architecture

- MicroBlaze is implemented with a **Harvard memory architecture**: instruction and data accesses are done in separate address spaces.
 - The **instruction address space** has a 32-bit virtual address range with 32-bit MicroBlaze (that is, handles up to 4GB of instructions), and can be extended up to a 64-bit physical address range.
 - The **data address space** has a default 32-bit range, and can be extended up to a 64-bit range (that is, handles from 4GB to 16GB of data).
- The instruction and data memory ranges can be made to overlap by mapping them both to the same physical memory. The latter is necessary for software debugging.
- Both instruction and data interfaces of MicroBlaze are default 32 bits wide and use big endian or little endian, bit-reversed format, depending on the selected endianness. MicroBlaze supports word, halfword, and byte accesses to data memory.
- Data accesses must be aligned (word accesses must be on word boundaries, halfword on halfword boundaries), unless the processor is configured to support unaligned exceptions. All instruction accesses must be word aligned.

Final Remarks

At the end of this lecture you should be able to:

- have a generic idea of MicroBlaze processor

To do:

- Complete the lab. 4