Designing a Custom AXI-Stream Peripheral

LECTURE 10

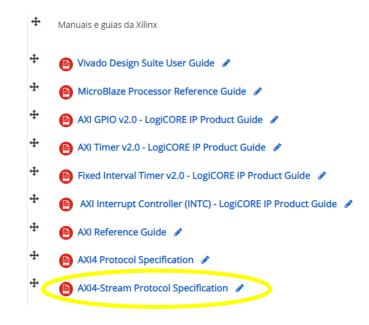
PEDRO SANTOS (SLIDES BY IOULIIA SKLIAROVA)

AXI4-Stream

- For high-speed streaming data in point-to-point communications.
- AXI4-Stream removes the requirement for an address phase altogether and allows unlimited data burst size.
- AXI4-Stream interfaces and transfers do not have address phases.
 - Therefore, AXI4-Stream is not considered to be memory-mapped.
- The AXI4-Stream protocol defines a single unidirectional channel for transmission of streaming data (with a handshaking data flow).
- The AXI4-Stream channel models the write data channel of AXI4.
 - No address channels nor read channels as in AXI4/AXI4-Lite.
 - Master (generates data) Slave (receives data) communication paradigm
- Use the AXI4-Stream protocol for applications that typically focus on a datacentric and data-flow paradigm where the concept of an address is not present or not required.

AXI4-Stream Interface Signals

Signal	Source	Description
ACLK	Clock source	The global clock signal. All signals are sampled on the rising edge of ACLK.
ARESETn	Reset source	The global reset signal. ARESETn is active-LOW.
TVALID	Master	TVALID indicates that the master is driving a valid transfer. A transfer takes place when both TVALID and TREADY are asserted.
TREADY	Slave	TREADY indicates that the slave can accept a transfer in the current cycle.
TDATA[(8n-1):0]	Master	TDATA is the primary payload that is used to provide the data that is passing across the interface. The width of the data payload is an integer number of bytes.
TSTRB[(n-1):0]	Master	TSTRB is the byte qualifier that indicates whether the content of the associated byte of TDATA is processed as a data byte or a position byte.
TKEEP[(n-1):0]	Master	TKEEP is the byte qualifier that indicates whether the content of the associated byte of TDATA is processed as part of the data stream. Associated bytes that have the TKEEP byte qualifier deasserted are null bytes and can be removed from the data stream.
TLAST	Master	TLAST indicates the boundary of a packet.
TID[(i-1):0]	Master	TID is the data stream identifier that indicates different streams of data.
TDEST[(d-1):0]	Master	TDEST provides routing information for the data stream.
TUSER[(u-1):0]	Master	TUSER is user defined sideband information that can be transmitted alongside the data stream.

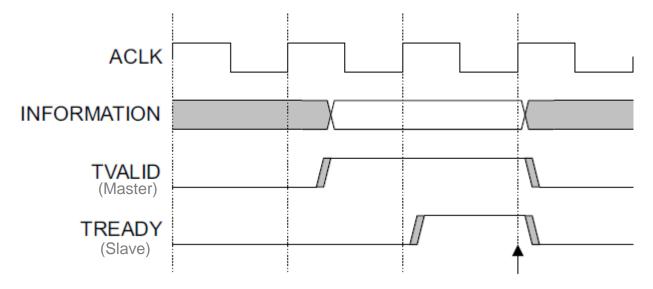


AXI4-Stream Handshake Process

- The TVALID and TREADY handshake determines when information is passed across the interface.
- A two-way flow control mechanism enables both the master and slave to control the rate at which the data and control information is transmitted across the interface.
- For a transfer to occur both the **TVALID** and **TREADY** signals must be asserted.
- Either TVALID or TREADY can be asserted first or both can be asserted in the same ACLK cycle.
- A master is not permitted to wait until **TREADY** is asserted before asserting **TVALID**. Once **TVALID** is asserted, it must remain asserted until the handshake occurs.
- A slave is permitted to wait for TVALID to be asserted before asserting the corresponding TREADY. If a slave asserts TREADY, it is permitted to deassert TREADY before TVALID is asserted.

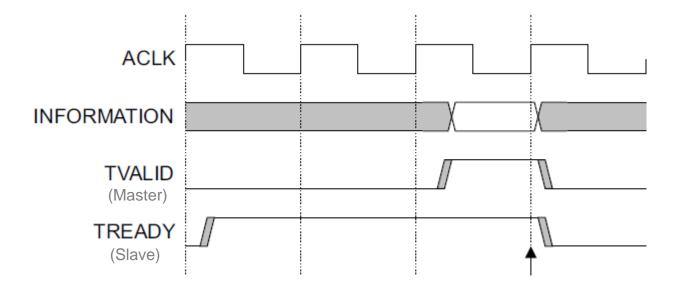
TVALID Before TREADY Handshake

- The master presents the data and control information and asserts the TVALID signal HIGH.
- Once the master has asserted TVALID, the data or control information from the master must remain unchanged until the slave drives the TREADY signal HIGH, indicating that it can accept the data and control information.
- In this case, transfer takes place once the slave asserts **TREADY** HIGH. The arrow shows when the transfer occurs.



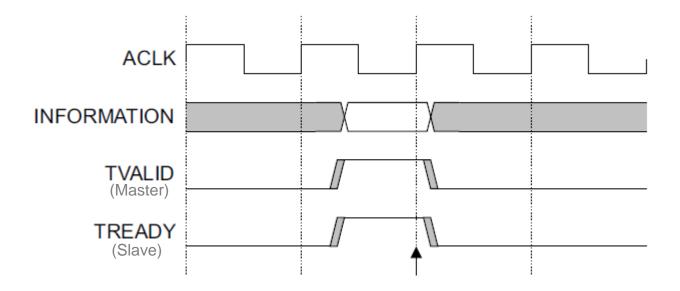
TREADY Before TVALID Handshake

- The slave drives TREADY HIGH before the data and control information is valid.
 This indicates that the destination can accept the data and control information in a single cycle of ACLK.
- In this case, transfer takes place once the master asserts TVALID HIGH. The arrow shows when the transfer occurs.



TVALID With TREADY Handshake

 The master asserts TVALID HIGH and the slave asserts TREADY HIGH in the same cycle of ACLK. In this case, transfer takes place in the same cycle as shown by the arrow.



Software use of coprocessor

```
//MicroBlaze Processor Fast Simplex Link (FSL) Interface Macros
//The FSL channels are dedicated unidirectional point-to-point data streaming interfaces.

void ReverseEndiannessHw(int* pDst, int* pSrc, unsigned int size)
{
   int* p;

   for (p = pSrc; p < pSrc + size; p++, pDst++)
   {
      putfslx(*p, 0, FSL_DEFAULT);
      //macros:
      // *p - value to send,
      // 0 - FSL (interface) identifier: a literal in the range of 0 to 7 (0 to 15 for MicroBlaze
      // v7.00.a and later).
      // FSL_DEFAULT - flags (by default)

      getfslx(*pDst, 0, FSL_DEFAULT);
   }
}</pre>
```

Example 1/Lab.8 – Part 1

REVERSE ENDIANNESS

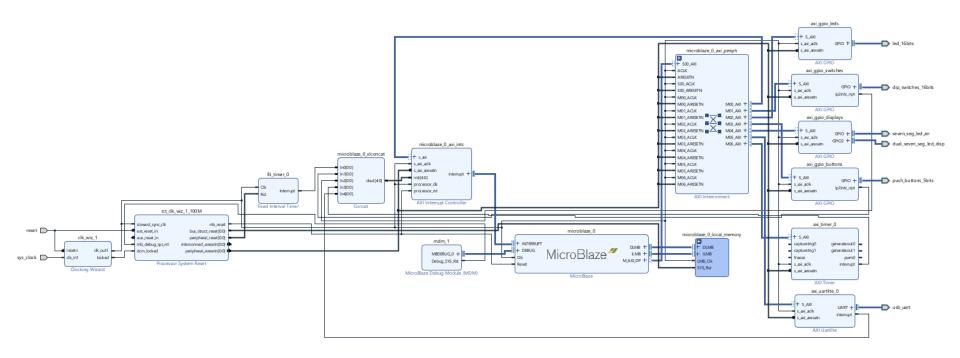
Example 1 / Lab.8-Part 1

Reverse endianness

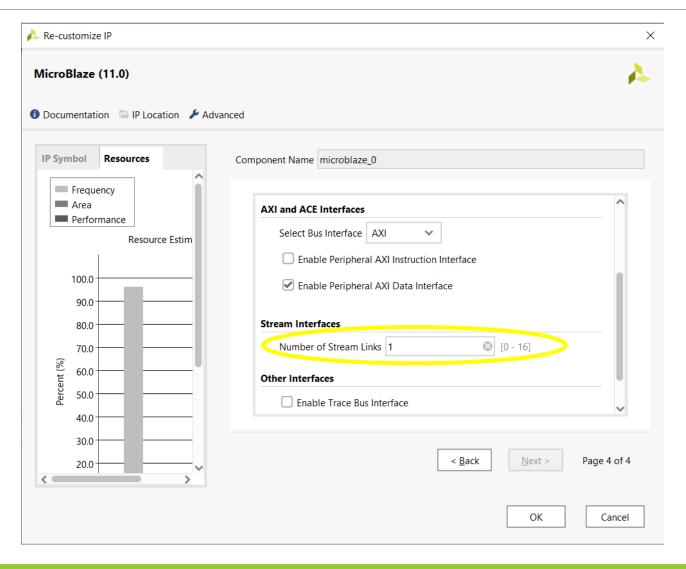
- Endianness is the order of bytes in a word of digital data.
- Endianness is primarily expressed as big-endian or little-endian.
- A big-endian system stores the most significant byte of a word at the smallest memory address and the least significant byte at the largest.
- A little-endian system, in contrast, stores the least-significant byte at the smallest address.
- 0xAB347801 => 0x017834AB

Population count (Hamming weight)

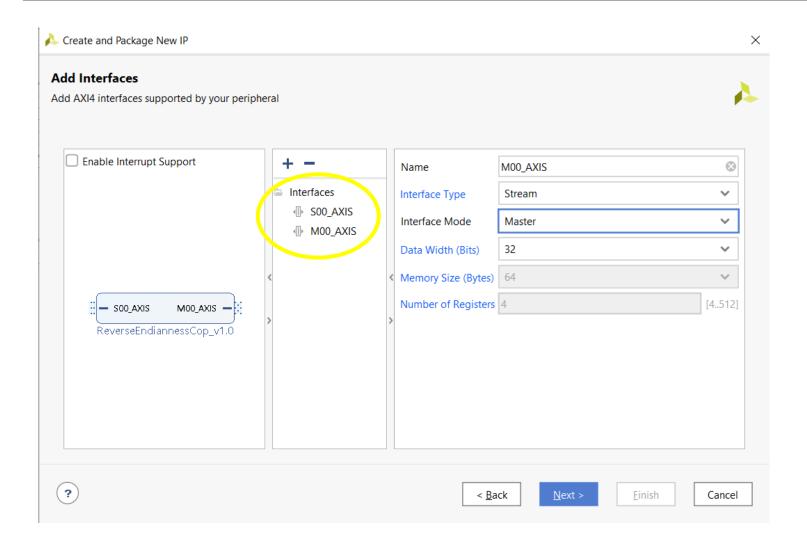
Example 1 – Starting Point



Adding Stream Links to MicroBlaze



Create & Package New IP



Code Description 1

generic (

M_AXIS_TSTRB M AXIS TLAST

M AXIS TREADY

reversedData

readEnabled

dataValid

);

port (

C S AXIS TDATA WIDTH

S_AXIS_ACLK : in std_logic;
S AXIS ARESETN : in std logic;

architecture arch_imp of ReverseEndiannessCop_v1_0 is component ReverseEndiannessCop_v1_0_S00_AXIS is

: out std logic; S AXIS TREADY : in std logic_vector(C_S_AXIS_TDATA_WIDTH-1 downto 0); S AXIS TDATA : in std logic vector((C S AXIS TDATA WIDTH/8)-1 downto 0); S AXIS TSTRB S AXIS TLAST : in std logic; S AXIS TVALID : in std logic; dataValid : out std logic; reversedData : out std logic vector(C S AXIS TDATA WIDTH-1 downto 0); : in std logic readEnabled end component ReverseEndiannessCop v1 0 S00 AXIS; component ReverseEndiannessCop v1 0 M00 AXIS is generic (C M AXIS TDATA WIDTH : integer := 32; C M START COUNT : integer := 32); port (M AXIS ACLK : in std logic; M AXIS ARESETN : in std logic; M AXIS TVALID : out std logic; M AXIS TDATA : out std logic vector(C M AXIS TDATA WIDTH-1 downto 0);

: out std logic;

: in std logic;

: in std logic;

: out std logic

end component ReverseEndiannessCop v1 0 M00 AXIS;

: out std logic vector((C M AXIS TDATA WIDTH/8)-1 downto 0);

: in std logic vector(C M AXIS TDATA WIDTH-1 downto 0);

: integer := 32

ReverseEndiannessCop_v1_0

Code Description 2

ReverseEndiannessCop _v1_0_S00_AXIS.vhd

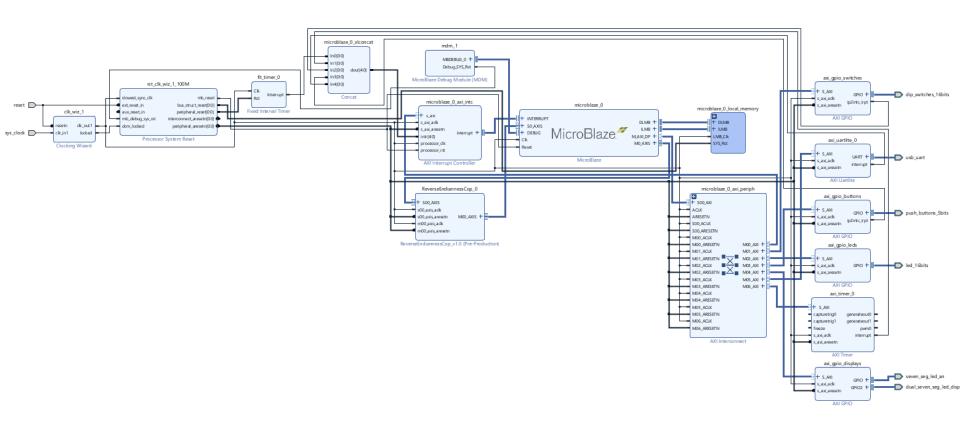
- Master: MicroBlaze (Master-MB)
- Slave: RevEnd module (Slave-Module)

```
architecture arch_imp of ReverseEndiannessCop_v1_0_S00_AXIS is
   signal s ready
                 : std logic;
   signal s validOut : std logic;
   signal s dataOut : std logic vector(C S AXIS TDATA WIDTH-1 downto 0);
   1) there's no data to sent by Master-Module or,
   process(S AXIS ACLK)
                                                         2) Master-MB is reading
   begin
      if (rising edge (S AXIS ACLK)) then
          if (S_AXIS_ARESETN = '0') then
             s validOut <= '0';
             s dataOut <= (others => '0');
          elsif (S AXIS TVALID = '1') then _____ If Master-MB reports valid data to transmit (TVALID high)...
             if (s_ready = '1') then ______ If Slave-Module is ready to receive...
                 s validOut <= '1';
                                                                                   Read data, perform computation,
                 and set Master-Module TVALID high
                             S AXIS TDATA(23 downto 16) & S AXIS TDATA(31 downto 24);
                                                                                   (signal that results are valid)
             end if;
          elsif (readEnabled = '1') then If Master-MB is reading, Master-Module can have
             s validOut <= '0';</pre>
                                                      TVALID signal set to low (no longer necessary)
          end if:
       end if:
   end process;
   dataValid
               <= s validOut;
   reversedData <= s dataOut;</pre>
   S AXIS TREADY <= s ready;
end arch imp;
```

ReverseEndiannessCop _v1_0_M00_AXIS.vhd

- Master: RevEnd module (Master-Module)
- Slave: MicroBlaze (Slave-MB)

Example 1 Block Design



Lab.8-Part 1 — Reverse Endianness

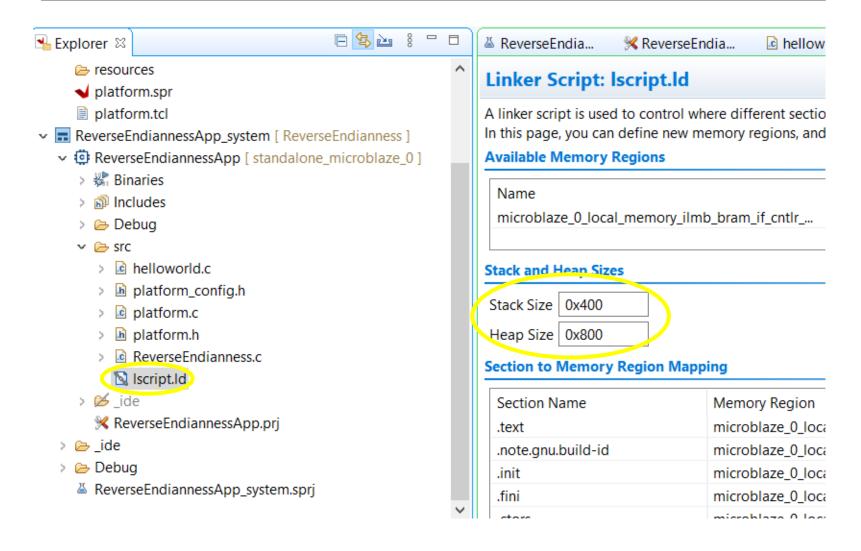
- 1. Import Block Design
- 2. Configure the MicroBlaze to have one pair of stream links
- 3. Create and Package new IP (*ReverseEndiannessCop*)
- 4. Add IP
- 5. Edit in IP Packager (the code is given on eLearning)
- 6. Generate output products
- 7. Create HDL Wrapper
- Generate Bitstream
 - o set_property CONFIG_VOLTAGE 3.3 [get_designs synth_1]
 - o set_property CFGBVS VCCO [get_designs synth_1]
- 9. Export Hardware
- 10. Launch Vitis

Sw Project (Vitis), Stack Size & Performance Eval.

Vitis

- The C code is given on eLearning ReverseEndianness.c
- There is no need to correct the IP makefiles.
- The code processes an array of N (N=4000) integers (32b = 4B)
- The code uses the AXI timer to measure time
- The code does not work. Why?
- #define N 4000
- int srcData[N], dstData[N];
- What is the size of data?
- Where do these data reside?

Vitis – Changing the Stack Size



Performance Evaluation

```
int main()
                                                                      ReverseEndianness.c
   int srcData[N], dstData[N];
   unsigned int timeElapsed;
   init platform();
   xil_printf("\n\rSoftware Only vs. Hardware Assisted Reverse Endianess Demonstration\n\r");
   RestartPerformanceTimer();
   srand(0);
   for (int i = 0; i < N; i++)
       srcData[i] = rand();
   timeElapsed = StopAndGetPerformanceTimer();
   xil printf("\n\rArray initialization time: %d microseconds\n\r",
              timeElapsed / (XPAR_CPU_M_AXI_DP_FREQ_HZ / 1000000)); // each period is 10 ns = 0.01 us
   PrintDataArray(srcData, min(8, N));
   xil printf("\n\r");
   // Software only
   RestartPerformanceTimer();
   ReverseEndiannessSw(dstData, srcData, N):
   timeElapsed = StopAndGetPerformanceTimer();
   xil printf("\n\rSoftware only reverse endianness time: %d microseconds",

    Software implementation

              timeElapsed / (XPAR CPU M AXI DP FREO HZ / 1000000));
   PrintDataArray(dstData, min(8, N));
   xil printf("\n\rChecking result: %s\n\r",
              CheckReversedEndianness(srcData, dstData, N) ? "OK" : "Error");
   // Hardware assisted
   RestartPerformanceTimer();
   ReverseEndiannessHw(dstData, srcData, N);
   timeElapsed = StopAndGetPerformanceTimer();
   xil printf("\n\rHardware assisted reverse endianness time: %d microseconds",
              timeElapsed / (XPAR_CPU_M_AXI_DP_FREQ_HZ / 1000000));
                                                                                    Hardware implementation
   PrintDataArray(dstData, min(8, N));
   xil_printf("\n\rChecking result: %s\n\r",
              CheckReversedEndianness(srcData, dstData, N) ? "OK" : "Error");
     cleanup_platform();
   return 0;
```

Compare time that elapses with SW and HW (co-processor) implementations.

Example 2 / Lab.8 – Part 2

POPULATION COUNT

Lab.8-Part 2 — Pop Count

Reverse endianness

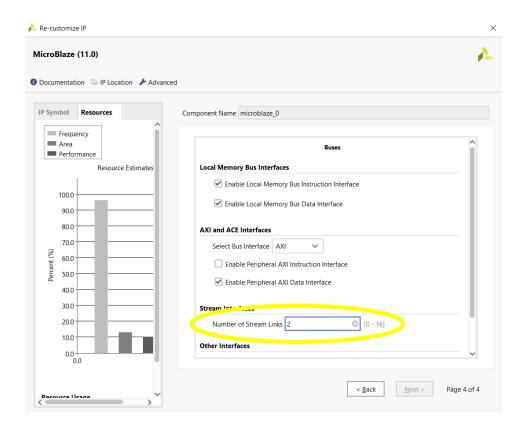
Population count (Hamming weight)

- the number of non-zero entries ('1' bits) in a word of data.
- o 0xAB347801 => 10101011_00110100_01111000_00000001 => 13

Lab.8 – **Part 2** is expected for submission. Submission form to open shortly in Elearning.

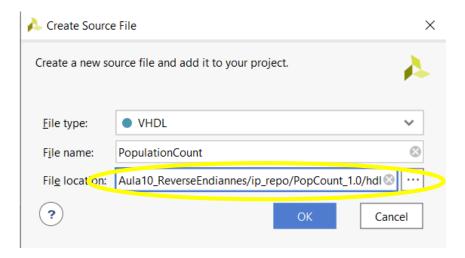
Lab.8-Part 2 — Starting Point

- Continue to work on the same project
- Change the number of stream links in the MicroBlaze to 2



Example 2 – Add New IP

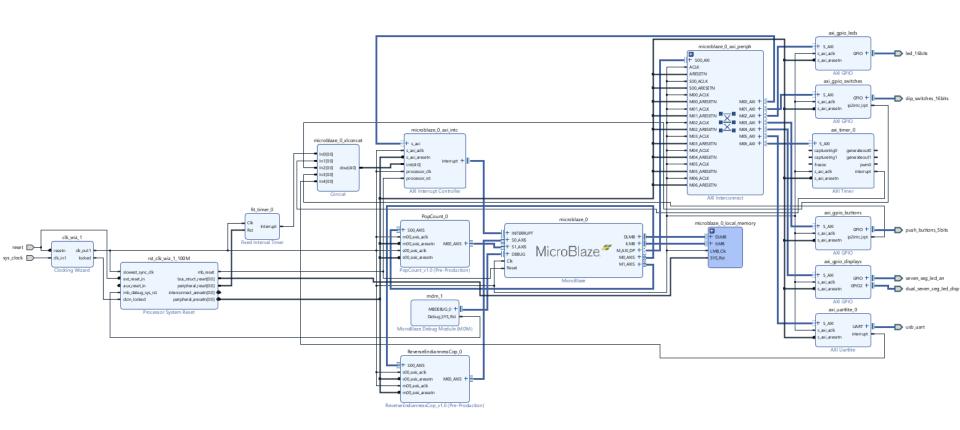
- Create and Package new IP PopCount
- Edit in IP packager
- Change the three "default" files like in Example 1
- Create a new source PopulationCount.vhd (the code is given on eLearning)



• Instantiate the PopulationCount module in the slave stream interface:

```
calc: PopulationCount
  generic map(N => C_S_AXIS_TDATA_WIDTH)
  port map( dataIn => ...);
```

Example 2 – Block Design



Example 2 – Further Steps

- Hw project (VIVADO)
 - 1. Generate output products
 - 2. Create HDL Wrapper
 - 3. Generate Bitstream
 - 4. Export Hardware
 - 5. Launch Vitis
- Sw project (VITIS)
 - 1. Write the C code (on the basis of the ReverseEndianness example)
 - 2. There is no need to correct the IP makefiles
 - 3. Configure the right stack size

Sw Project & For Loop

For Loop VHDL Statement

- A for loop statement is a sequential statement that can be used inside a process.
- A for loop includes a specification of how many times the body of the loop is to be executed:

Note: It is unnecessary to declare the identifier entity.

- The for loop statement is used whenever an operation needs to be repeated.
- The loop is unrolled statically the number of loop iterations must be known at compile time.
- Loop unrolling is a systematic method of achieving parallelism that can be automated.
- This comes at a cost of a larger fabric footprint (more FPGA area).

Population Count With a For Loop

```
entity PopulationCount is
 generic(N : positive := 4);
 port(dataIn : in std logic vector(N-1 downto 0);
       cntOut : out std logic vector(N-1 downto 0));
end PopulationCount;
architecture Behavioral of PopulationCount is
    signal s cnt : natural range 0 to N;
begin
    process (dataIn)
        variable v cnt : natural range 0 to N;
    begin
        v cnt := 0;
        for i in 0 to N-1 loop
            if dataIn(i) = '1' then
                v cnt := v cnt + 1;
            end if:
        end loop;
        s cnt <= v cnt;
    end process;
    cntOut <= std logic vector(to_unsigned(s cnt, N));</pre>
end Behavioral;
```

A long sequence of 31 adders will be generated.

Variables

- For loops are often used with variables.
- Variables are declared in the declaration part of processes:

```
variable_declaration 
variable identifier { , ... } : subtype_indication
[:= expression] ;
```

• The syntax of a variable assignment statement is given by the rule

```
variable_assignment_statement 
[label :] name := expression ;
```

 A variable assignment immediately overwrites the variable with a new value (a signal assignment, on the other hand, schedules a new value to be applied to a signal at some later time).

For Loop vs For Generate

- The for loops are sequential statements, containing sequential statements (i.e. each iteration is sequenced to be executed after the previous one).
- The for-generate loops are concurrent statements, containing
 concurrent statements, and this is how you can use it to make several
 instances of a component.
- For-generate loops are used when specifying the exact hardware structure.
- For loops are more suited for behavioral descriptions.

Final Remarks

At the end of this lecture you should be able to:

- Design custom hardware modules interacting with the MicroBlaze through AXI-Stream interface
- Write C programs that make use of stream-connected custom hardware

To do:

- Construct the considered hardware platforms
- Test the given application in Vitis
- Complete Lab. 8 Part 2